

Promises are meant to be broken

Strictness Analysis for R

Author 1^{*}
Position1
Department1
Institution1
Street1 Address1
City1, State1 Post-Code1, Country1
first1.last1@inst1.edu

Author 2[†]
Position2a
Department2a
Institution2a
Street2a Address2a
City2a, State2a Post-Code2a, Country2a
first2.last2@inst2a.com
Position2b
Department2b
Institution2b
Street3b Address2b
City2b, State2b Post-Code2b, Country2b
first2.last2@inst2b.org

Abstract

R evaluates function arguments lazily. When a function is called the formal parameters are bound to “promises” Arguments are stored in “promises” which are “forced” when the corresponding parameter is used in an evaluation context. A promise also contains a reference to the environment in which the argument has to be evaluated and a slot to cache the result of evaluation. This mechanism can potentially improve speed of R programs by preventing the evaluation of arguments that are not used. However, benchmarks obtained by instrumenting the R interpreter indicate that lazy evaluation is futile for majority of R programs and detrimental for their performance. R has a functional core and even most the basic control structures are implemented as functions. This results in the creation of a huge number of promises which get forced immediately but continue storing the unevaluated arguments until they are garbage collected. Furthermore, accessing the argument’s value requires the promise to look up the cache which introduces a level of indirection, increasing the possibility of cache misses.

In this paper, we present a static analysis technique on an intermediate representation of R code that identifies the promises that are forced and the order in which they are forced. This information can then be used by an optimizer to prevent the creation of promises and evaluate arguments in the correct order, ensuring unchanged program behavior. We validate the results of our static analysis from dynamic runs on a number of widely used R libraries.

^{*}with author1 note

[†]with author2 note

with paper note.

PL’17, New York, NY, USA

2017. 978-x-xxxx-xxxx-x/YY/MM...\$15.00

DOI: 10.1145/nnnnnnn.nnnnnnn




Figure 1. Structure of a promise.

CCS Concepts •Theory of computation → Program analysis; •Software and its engineering → Procedures, functions and subroutines; Multiparadigm languages;

Keywords R, Strictness analysis, Control-Flow analysis, Static analysis, Program analysis, Performance optimization

ACM Reference format:

Author 1 and Author 2. 2017. Promises are meant to be broken. In *Proceedings of ACM SIGPLAN Conference on Programming Languages*, New York, NY, USA, January 01–03, 2017 (PL’17), 2 pages. DOI: 10.1145/nnnnnnn.nnnnnnn

1 Introduction

Text of paper ...

2 Overview

2.1 Promise

A promise object has three slots; value, expression and environment. When a function is called, the actual arguments are wrapped inside promises. The promise also contains a reference to the

2.2 Background on R/Strictness

2.3 Example/IR representation

```
1 f1 <- function (x = 2, y = 3) {  
2   print (x)
```

```
3 f2 ( z = { x <- "Hello" })  
4 print ( x )  
5 }
```

3 Algorithm

4 Implementation

5 Experimental Results

- Percentage of promises evaluated in the first arguments.
- Percentage of promises with side effects.
- Percentage of promises never evaluated.
- Percentage of promises escaping to other functions and being evaluated.
- Promises which are evaluated in different order in different tests.
- Compare this with the result of strictness analysis.

6 Related Work

7 Conclusions

Acknowledgments

This material is based upon work supported by the National Science Foundation under Grant No. nnnnnnnn and Grant No. mmmmmmmm. Any opinions, findings, and conclusions or recommendations expressed in this material are those of the author and do not necessarily reflect the views of the National Science Foundation.

References

A Appendix

Text of appendix ...